

# DSBDA Assignment 2b

**Perform the following operations using Python on the Air Quality and Heart Diseases data sets**

1. Data Cleaning
2. Data Integration
3. Error Correcting
4. Data Model Building

## Import Python Libraries

```
In [1]: import pandas as pd  
import numpy as np
```

## Reading the air\_quality dataset -

```
In [2]: df = pd.read_csv("air_quality.csv")
df
```

```
Out[2]:
```

	Unnamed: 0	Ozone	Solar.R	Wind	Temp	Month	Day	Humidity
0	1	41.0	190.0	7.4	67	5	1	High
1	2	36.0	118.0	8.0	72	5	2	High
2	3	12.0	149.0	12.6	74	5	3	Low
3	4	18.0	313.0	11.5	62	5	4	NaN
4	5	NaN	NaN	14.3	56	5	5	High
...	...	...	...	...	...	...	...	...
148	149	30.0	193.0	6.9	70	9	26	Low
149	150	NaN	145.0	13.2	77	9	27	NaN
150	151	14.0	191.0	14.3	75	9	28	High
151	152	18.0	131.0	8.0	76	9	29	Medium
152	153	20.0	223.0	11.5	68	9	30	Low

153 rows × 8 columns

```
In [3]: df.isnull().sum()
```

```
Out[3]: Unnamed: 0      0
Ozone      37
Solar.R     7
Wind        2
Temp        0
Month        0
Day          0
Humidity    15
dtype: int64
```

# Data Cleaning

Fill the null values with the mean and the mode accordingly -

*Removing unwanted column from the dataset*

```
In [8]: df = df.drop(['Unnamed: 0'], axis = 1)
df
```

```
Out[8]:
```

	Ozone	Solar.R	Wind	Temp	Month	Day	Humidity
0	41.0	190.0	7.4	67	5	1	High
1	36.0	118.0	8.0	72	5	2	High
2	12.0	149.0	12.6	74	5	3	Low
3	18.0	313.0	11.5	62	5	4	NaN
4	NaN	NaN	14.3	56	5	5	High
...	...	...	...	...	...	...	...
148	30.0	193.0	6.9	70	9	26	Low
149	NaN	145.0	13.2	77	9	27	NaN
150	14.0	191.0	14.3	75	9	28	High
151	18.0	131.0	8.0	76	9	29	Medium
152	20.0	223.0	11.5	68	9	30	Low

153 rows × 7 columns

```
In [9]: df.columns
```

```
Out[9]: Index(['Ozone', 'Solar.R', 'Wind', 'Temp', 'Month', 'Day', 'Humidity'], dtype='object')
```

```
In [10]: df.isnull().sum()
```

```
Out[10]: Ozone      37  
Solar.R    7  
Wind       2  
Temp       0  
Month      0  
Day        0  
Humidity   15  
dtype: int64
```

```
In [11]: df['Ozone'].fillna(df['Ozone'].mean(), inplace=True)  
df['Solar.R'].fillna(df['Solar.R'].mean(), inplace=True)  
df['Wind'].fillna(df['Wind'].mean(), inplace=True)
```

Now we have to use mode in the case of Humidity because it is categorical dataset and we need to convert it in numerical form.

```
In [15]: df['Humidity'] = df['Humidity'].fillna(df['Humidity'].mode()[0])
```

```
In [16]: df
```

```
Out[16]:
```

	Ozone	Solar.R	Wind	Temp	Month	Day	Humidity
0	41.00000	190.000000	7.4	67	5	1	High
1	36.00000	118.000000	8.0	72	5	2	High
2	12.00000	149.000000	12.6	74	5	3	Low
3	18.00000	313.000000	11.5	62	5	4	High
4	42.12931	185.931507	14.3	56	5	5	High
...	...	...	...	...	...	...	...
148	30.00000	193.000000	6.9	70	9	26	Low
149	42.12931	145.000000	13.2	77	9	27	High
150	14.00000	191.000000	14.3	75	9	28	High
151	18.00000	131.000000	8.0	76	9	29	Medium
152	20.00000	223.000000	11.5	68	9	30	Low

153 rows × 7 columns

## Error Correcting

### Look for missing values

```
In [17]: df.isnull().sum()
```

```
Out[17]: Ozone      0
Solar.R    0
Wind       0
Temp       0
Month      0
Day        0
Humidity    0
dtype: int64
```

```
In [18]: df.describe()
```

```
Out[18]:
```

	Ozone	Solar.R	Wind	Temp	Month	Day
<b>count</b>	153.000000	153.000000	153.000000	153.000000	153.000000	153.000000
<b>mean</b>	42.129310	185.931507	9.945033	77.882353	6.993464	15.803922
<b>std</b>	28.693372	87.960267	3.520648	9.465270	1.416522	8.864520
<b>min</b>	1.000000	7.000000	1.700000	56.000000	5.000000	1.000000
<b>25%</b>	21.000000	120.000000	7.400000	72.000000	6.000000	8.000000
<b>50%</b>	42.129310	194.000000	9.700000	79.000000	7.000000	16.000000
<b>75%</b>	46.000000	256.000000	11.500000	85.000000	8.000000	23.000000
<b>max</b>	168.000000	334.000000	20.700000	97.000000	9.000000	31.000000

```
In [21]: df.dtypes
```

```
Out[21]: Ozone      float64
Solar.R    float64
Wind       float64
Temp       int64
Month      int64
Day        int64
Humidity   object
dtype: object
```

## Data Transformation

```
In [22]: from sklearn.preprocessing import LabelEncoder
lb = LabelEncoder()
```

```
In [23]: df["Humidity"] = lb.fit_transform(df["Humidity"])
df
```

```
Out[23]:
```

	Ozone	Solar.R	Wind	Temp	Month	Day	Humidity
0	41.00000	190.000000	7.4	67	5	1	0
1	36.00000	118.000000	8.0	72	5	2	0
2	12.00000	149.000000	12.6	74	5	3	1
3	18.00000	313.000000	11.5	62	5	4	0
4	42.12931	185.931507	14.3	56	5	5	0
...	...	...	...	...	...	...	...
148	30.00000	193.000000	6.9	70	9	26	1
149	42.12931	145.000000	13.2	77	9	27	0
150	14.00000	191.000000	14.3	75	9	28	0
151	18.00000	131.000000	8.0	76	9	29	2
152	20.00000	223.000000	11.5	68	9	30	1

153 rows × 7 columns

## Data Model Building

### Splitting x and y

```
In [24]: x = df.iloc[:, [0, 3]].values
y = df['Humidity']
```

### Splitting the training and testing data

```
In [25]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25)
```

### Importing logistic regression

```
In [26]: from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(x_train, y_train)
```

```
Out[26]: LogisticRegression()
```

```
In [27]: pred = model.predict(x_test);
pred
```

```
Out[27]: array([0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 2, 0, 0,
               0, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 2, 0, 0, 0, 0])
```



```
In [29]: from sklearn.metrics import classification_report
print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
0	0.38	0.81	0.52	16
1	0.00	0.00	0.00	13
2	0.20	0.10	0.13	10
accuracy			0.36	39
macro avg	0.19	0.30	0.22	39
weighted avg	0.21	0.36	0.25	39

C:\Users\Rishabh\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero\_division' parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\Users\Rishabh\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero\_division' parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\Users\Rishabh\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero\_division' parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
In [ ]:
```